
Storage Management, Protection and Security

Module 7

Dr. Naveenkumar J
Associate Professor,
PRP- 217 - 4

Hard Disk Drive

- A Hard Disk Drive (HDD) is a **non-volatile data storage device** commonly found in computers, laptops, servers, and many electronic devices.

- It **uses magnetic storage to store and retrieve digital information by spinning multiple rigid disks (platters) coated with magnetic material.**

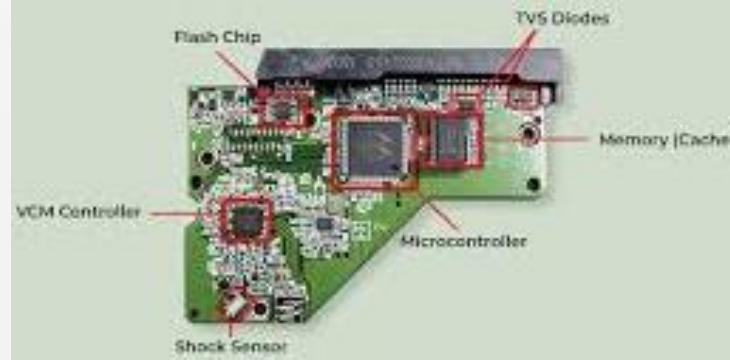
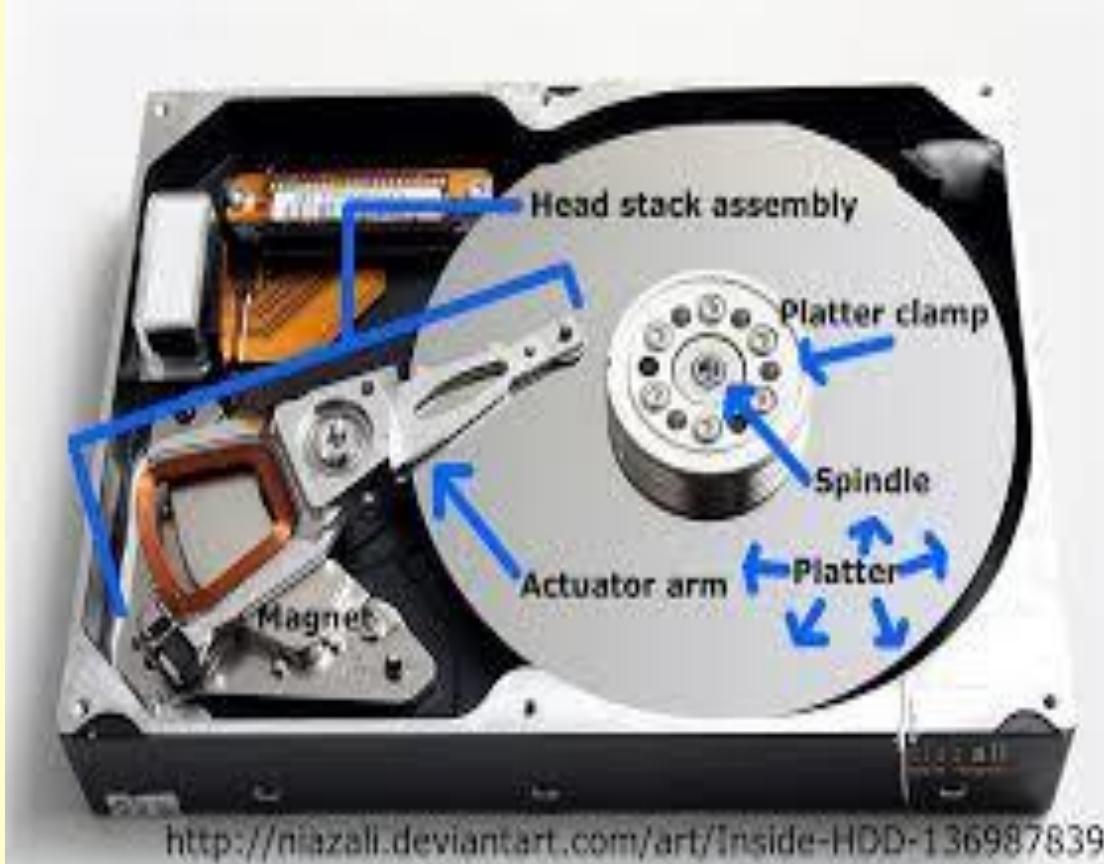
HDD Key Characteristics

- **Function:** Stores operating systems, applications, documents, media files, and virtually all user and system data.
- **Structure:** Contains spinning platters, read/write heads, actuator arm, controller electronics, and is housed in a protective casing.
- **Working Principle:** Data is encoded as tiny magnetized zones on the platters. As platters spin, read/write heads move very close to the surface to read or change the magnetization, representing binary data (0s and 1s).
- **Capacity:** Modern HDDs range from hundreds of gigabytes (GB) to multiple terabytes (TB)

HDD – Simple Working

- The platters spin at high speed (5400–15000 RPM).
- **Read/write heads** attached to arms quickly move to locate the right place on the platter.
- The **controller board manages all operations and communicates** with the computer via a connector (SATA, IDE, SCSI, etc.).
- Data stays intact even when the device is powered off (non-volatile).

HDD – Components



HDD – Components

❑ Platters

- ❑ Flat, circular disks (usually made of aluminum or glass, coated with magnetic material) where data is actually stored.
- ❑ A standard HDD contains multiple platters stacked vertically; each surface can store data.

❑ Spindle

- ❑ The rotating axis/wheel holding the platters in place.
- ❑ Spins the platters at high speeds (typically 5400, 7200, or 10,000+ RPM), enabling rapid data access.

HDD – Components

❑ **Read/Write Arm (Actuator Arm)**

- ❑ A thin, pivoting metallic arm that moves across the platters.
- ❑ Positions the read/write head over the correct track on the platter for reading or writing data.

❑ **Actuator (Voice Coil Motor)**

- ❑ The device (often an electromagnet) that controls the precise movement of the read/write arm.
- ❑ Allows rapid, accurate movement to the required data location.

HDD – Components

□ **Read/Write Head**

- Microscopic electromagnetic devices located at the tip of each actuator arm.
- Flies just above the platter surface, reading or altering the magnetic fields to store/retrieve bits.
- Each platter surface has its own dedicated head.

□ **Controller Board (PCB)**

- The electronic circuit board mounted on the underside of the drive.
- Manages all internal communication, power regulation, and data transfer between the HDD and host (computer).

HDD – Components

❑ **Connector Interface**

- ❑ External data and power connectors (SATA, PATA/IDE, SCSI, etc.) for connecting to the motherboard and power supply.

HDD – Components

| Component | Function/Role |
|------------------|--------------------------------------|
| Casing | Protection and physical support |
| Platters | Magnetic data storage |
| Spindle | Rotates platters for data access |
| Read/Write Arm | Positions the heads over data tracks |
| Actuator | Moves the arm precisely |
| Read/Write Head | Reads/writes magnetic data |
| Controller Board | Controls logic, communication, power |
| Connector | Connects power/data to computer |

Disk Scheduling Algorithm

- ❑ A disk scheduling algorithm is a **strategy used by operating systems** to **decide the order in which disk I/O** (input/output) **requests are serviced.**
- ❑ These algorithms are vital in multi-tasking systems or environments with heavy disk usage, because multiple read/write requests may arrive simultaneously or out of sequence.

Disk Scheduling Algorithm – Why important

- Minimizes overall waiting time for disk accesses.
- Reduces mechanical movement of the disk's read/write heads, decreasing seek time and improving performance.
- Prevents starvation and ensures fairness among requests.

Disk Scheduling Algorithm – Goals

- **Optimize:**
 - **Seek time** (time to move disk head between tracks)
 - **Rotational latency** (time for required sector to rotate under head)
 - **Throughput** (number of serviced requests per time unit)
 - **Fairness** (avoid neglecting distant requests)

Disk Scheduling Algorithm – Metrics

□ Seek Time

- The **time taken** for the **disk's read/write head to move from its current position to the track containing the desired data.**

□ Example

- Suppose the head is at track 20 and your data is at track 80. If each track jump takes 0.5 ms, total seek time = 60 tracks \times 0.5 ms = 30 ms.

Disk Scheduling Algorithm – Metrics

❑ Rotational Latency

❑ The time it takes for the desired disk sector to rotate under the read/write head after the head has reached the correct track. It depends on disk rotation speed.

❑ Example:

- ❑ A disk spins at 7200 RPM (Revolutions Per Minute). One full rotation takes $60\text{sec}/7200 = 0.0083\text{sec} = 8.3\text{ms}$
- ❑ The average rotational latency is half of this: 4.15 ms.

Disk Scheduling Algorithm – Metrics

❑ Transfer Time

- ❑ The time required to **actually transfer (read/write)** the data once the head is positioned correctly.

❑ Example:

If the disk can transfer data at 100 MB/s and you want to read

$$1 \text{ MB, Transfer Time} = \frac{1 \text{ MB}}{100 \text{ MB/s}} = 0.01 \text{ sec} = 10 \text{ ms.}$$

Disk Scheduling Algorithm – Metrics

❑ Throughput

- ❑ The number of disk requests serviced in a given period (requests per second).

❑ Example:

- ❑ If 70 I/O requests are completed in 10 seconds,

$$\text{Throughput} = \frac{70}{10} = 7 \text{ requests/sec.}$$

Disk Scheduling Algorithm – Metrics

❑ Fairness

- ❑ The extent to which requests—whether near or far from the current head position—are serviced without undue delay or starvation.

❑ Example:

- ❑ SSTF may be fast for short seeks, but is less fair (may starve far requests), while SCAN and C-SCAN offer balanced fairness.

Disk Scheduling Algorithms

| Algorithm | Definition |
|---------------------------------|--|
| FCFS (First-Come, First-Served) | Requests are processed in the exact order they arrive. Simple and fair, but not always efficient. |
| SSTF (Shortest Seek Time First) | The request closest to the current position of the disk arm is handled first, which minimizes immediate head movement. |
| SCAN (Elevator) | The disk arm moves in one direction, servicing all requests along its path. When it reaches the end, it reverses and services requests on the way back, like an elevator. |
| C-SCAN (Circular SCAN) | The disk arm moves in one direction, servicing requests until it reaches the end. Then, it quickly returns to the start (without servicing on the way back) and repeats. |
| LOOK | The disk arm moves only as far as the last request in each direction, then reverses, instead of going all the way to the physical end of the disk. |
| C-LOOK | The arm goes from the current position to the furthest request in one direction, then jumps to the furthest request in the other direction and continues, skipping over empty parts of the disk. |

Disk Scheduling Algorithm

- ❑ Suppose the head of moving disk with 200 tracks, numbered 0 to 199 is currently serving a request at track 132 and has just finished a request at track 120.
- ❑ If the queue of requests is kept in the FIFO order 38, 55, 86, 123, 147, 91, 177, 115, 94, 150, 100, 175 and 130, 185.
- ❑ What is total head movement to satisfy these request for the all-disk scheduling algorithms?

Disk Scheduling Algorithm

□ Given Data

- Tracks: 0 to 199
- Current Head Position: 132
- Previous Head Position: 120 (**Indicates the head is moving towards higher track numbers**)
- Request Queue (FIFO): 38, 55, 86, 123, 147, 91, 177, 115, 94, 150, 100, 175, 130, 185

Disk Scheduling Algorithm - FCFS

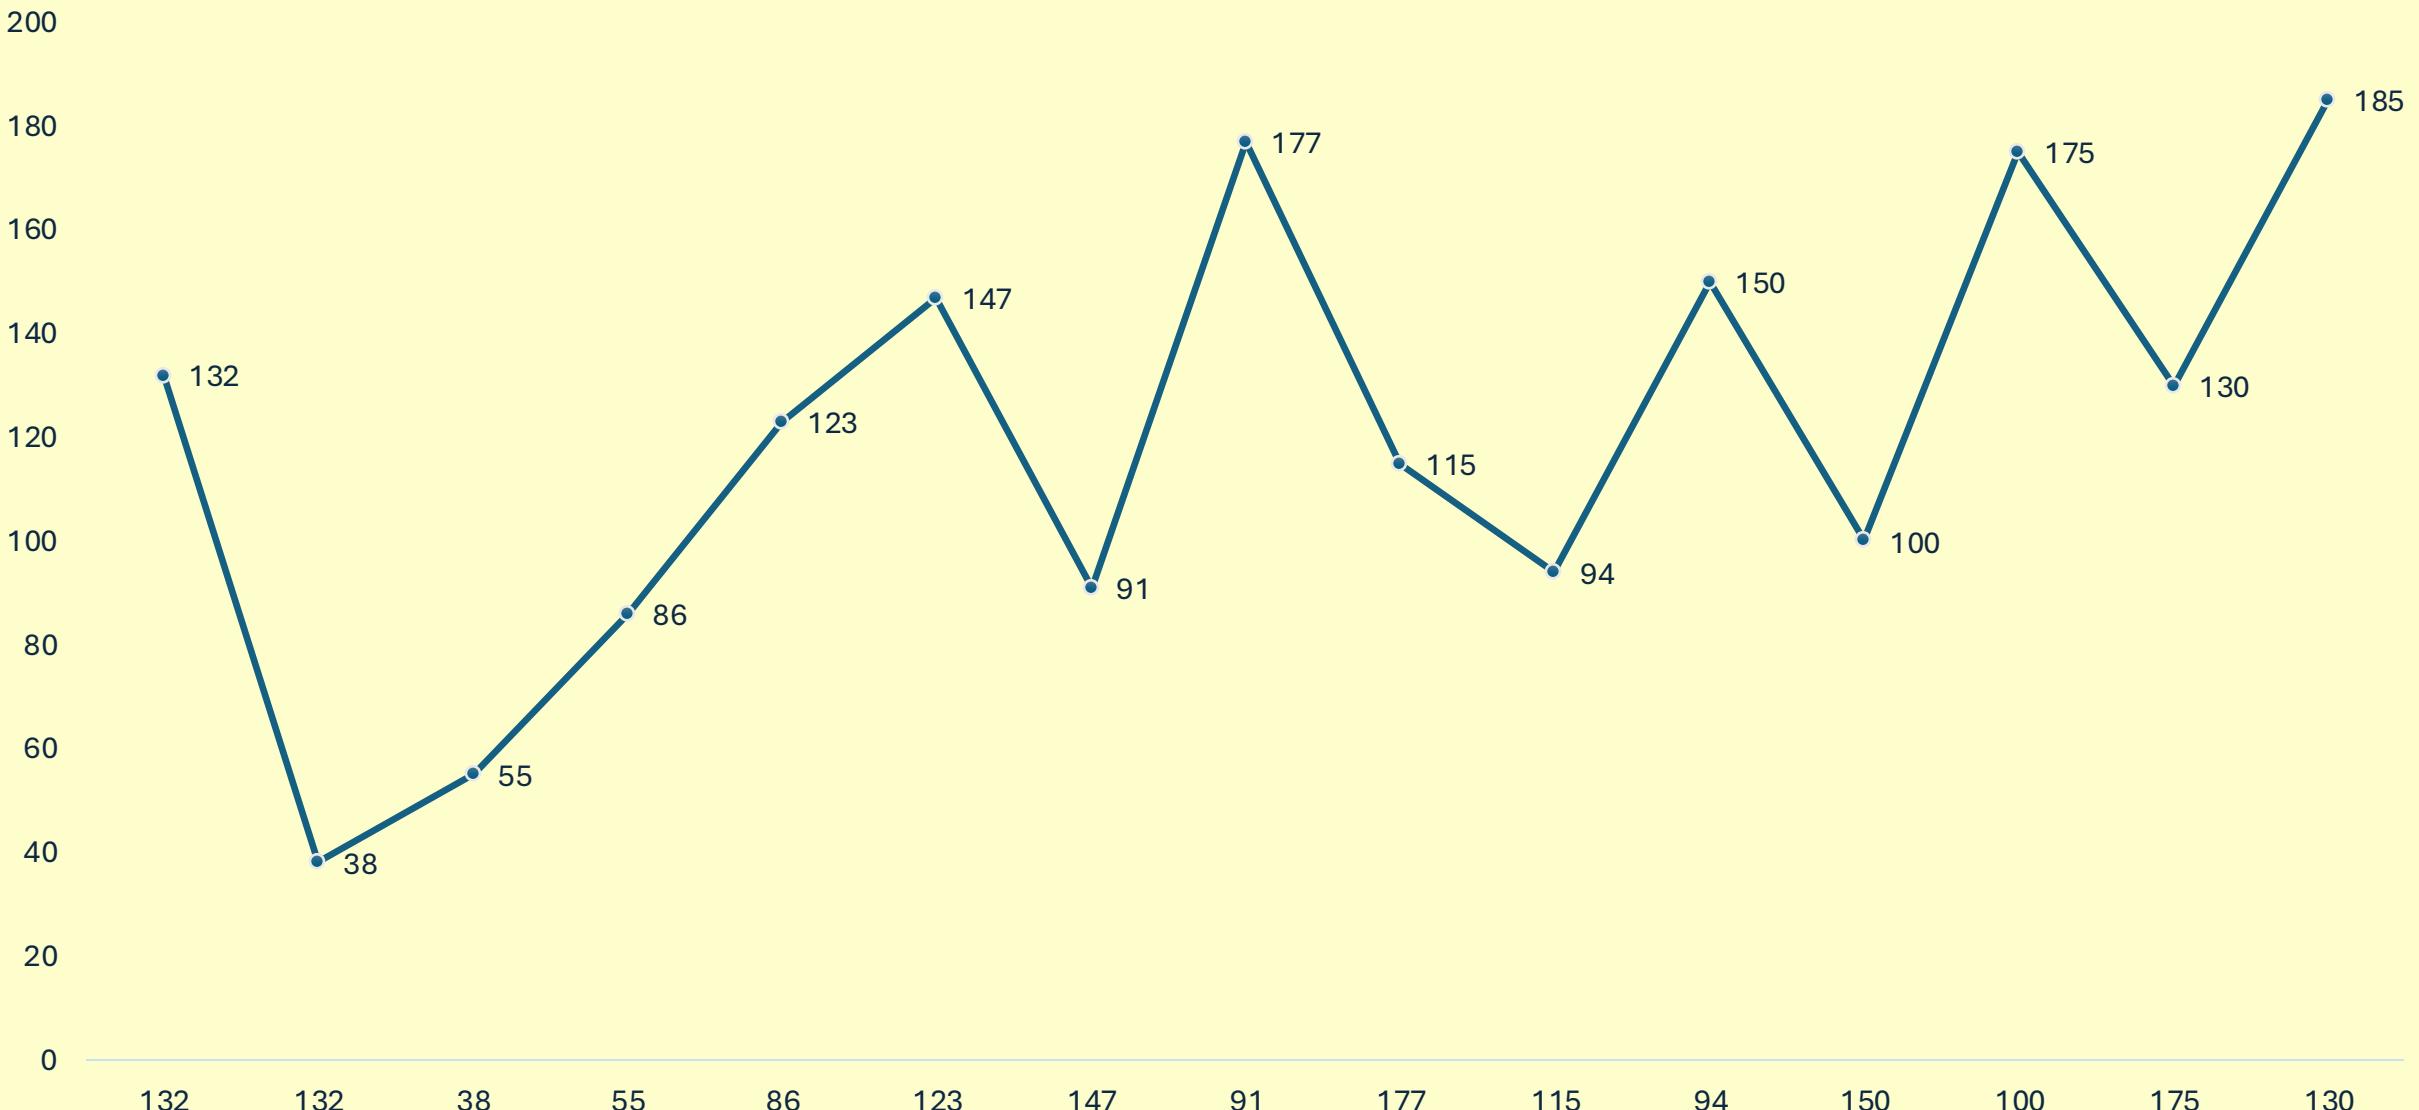
- ❑ In the FCFS algorithm, the requests are serviced in the order they appear in the queue.
- ❑ **Path:** **132 → 38 → 55 → 86 → 123 → 147 → 91 → 177 → 115 → 94 → 150 → 100 → 175 → 130 → 185**

Disk Scheduling Algorithm - FCFS

| Step | Current Head Position | Request Serviced | Head Movement | Total Movement |
|------|-----------------------|------------------|--------------------|----------------|
| 1 | 132 | 38 | $ 132 - 38 = 94$ | 94 |
| 2 | 38 | 55 | $ 38 - 55 = 17$ | 111 |
| 3 | 55 | 86 | $ 55 - 86 = 31$ | 142 |
| 4 | 86 | 123 | $ 86 - 123 = 37$ | 179 |
| 5 | 123 | 147 | $ 123 - 147 = 24$ | 203 |
| 6 | 147 | 91 | $ 147 - 91 = 56$ | 259 |
| 7 | 91 | 177 | $ 91 - 177 = 86$ | 345 |
| 8 | 177 | 115 | $ 177 - 115 = 62$ | 407 |
| 9 | 115 | 94 | $ 115 - 94 = 21$ | 428 |
| 10 | 94 | 150 | $ 94 - 150 = 56$ | 484 |
| 11 | 150 | 100 | $ 150 - 100 = 50$ | 534 |
| 12 | 100 | 175 | $ 100 - 175 = 75$ | 609 |
| 13 | 175 | 130 | $ 175 - 130 = 45$ | 654 |
| 14 | 130 | 185 | $ 130 - 185 = 55$ | 709 |

Disk Scheduling Algorithm - FCFS

—●—FCFS Algo

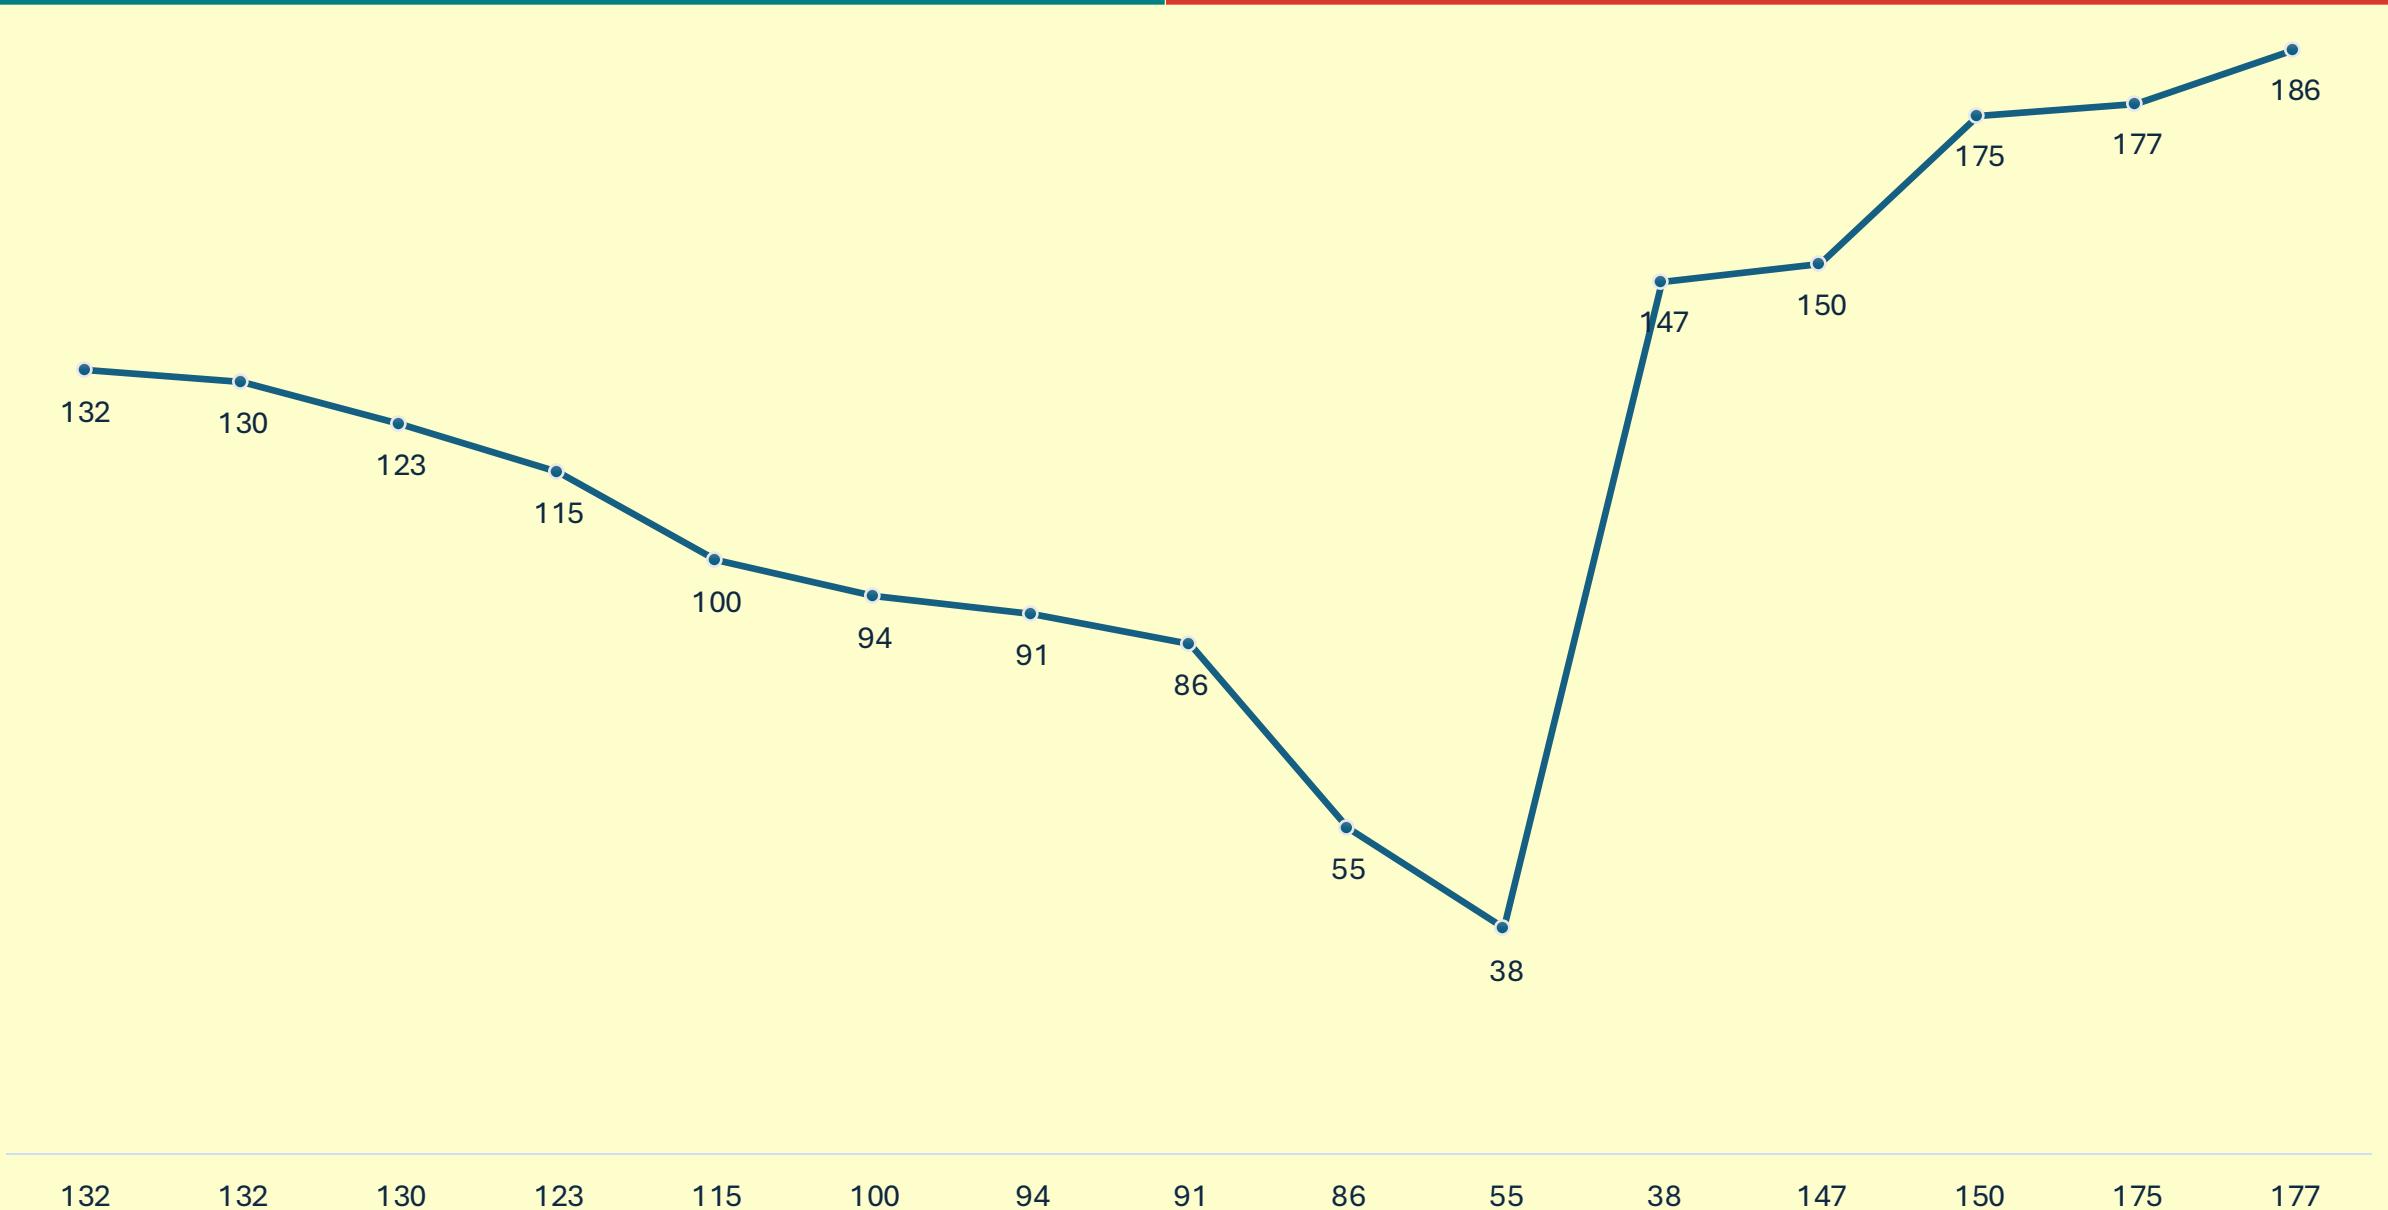


Disk Scheduling Algorithm - SSTF

| Step | Current Head Position | Request Serviced | Head Movement | Total Movement |
|------|-----------------------|------------------|--------------------|----------------|
| 1 | 132 | 130 | $ 132 - 130 = 2$ | 2 |
| 2 | 130 | 123 | $ 130 - 123 = 7$ | 9 |
| 3 | 123 | 115 | $ 123 - 115 = 8$ | 17 |
| 4 | 115 | 100 | $ 115 - 100 = 15$ | 32 |
| 5 | 100 | 94 | $ 100 - 94 = 6$ | 38 |
| 6 | 94 | 91 | $ 94 - 91 = 3$ | 41 |
| 7 | 91 | 86 | $ 91 - 86 = 5$ | 46 |
| 8 | 86 | 55 | $ 86 - 55 = 31$ | 77 |
| 9 | 55 | 38 | $ 55 - 38 = 17$ | 94 |
| 10 | 38 | 147 | $ 38 - 147 = 109$ | 203 |
| 11 | 147 | 150 | $ 147 - 150 = 3$ | 206 |
| 12 | 150 | 175 | $ 150 - 175 = 25$ | 231 |
| 13 | 175 | 177 | $ 175 - 177 = 2$ | 233 |
| 14 | 177 | 185 | $ 177 - 185 = 8$ | 241 |

The head moves to the closest pending request.

Disk Scheduling Algorithm - SSTF



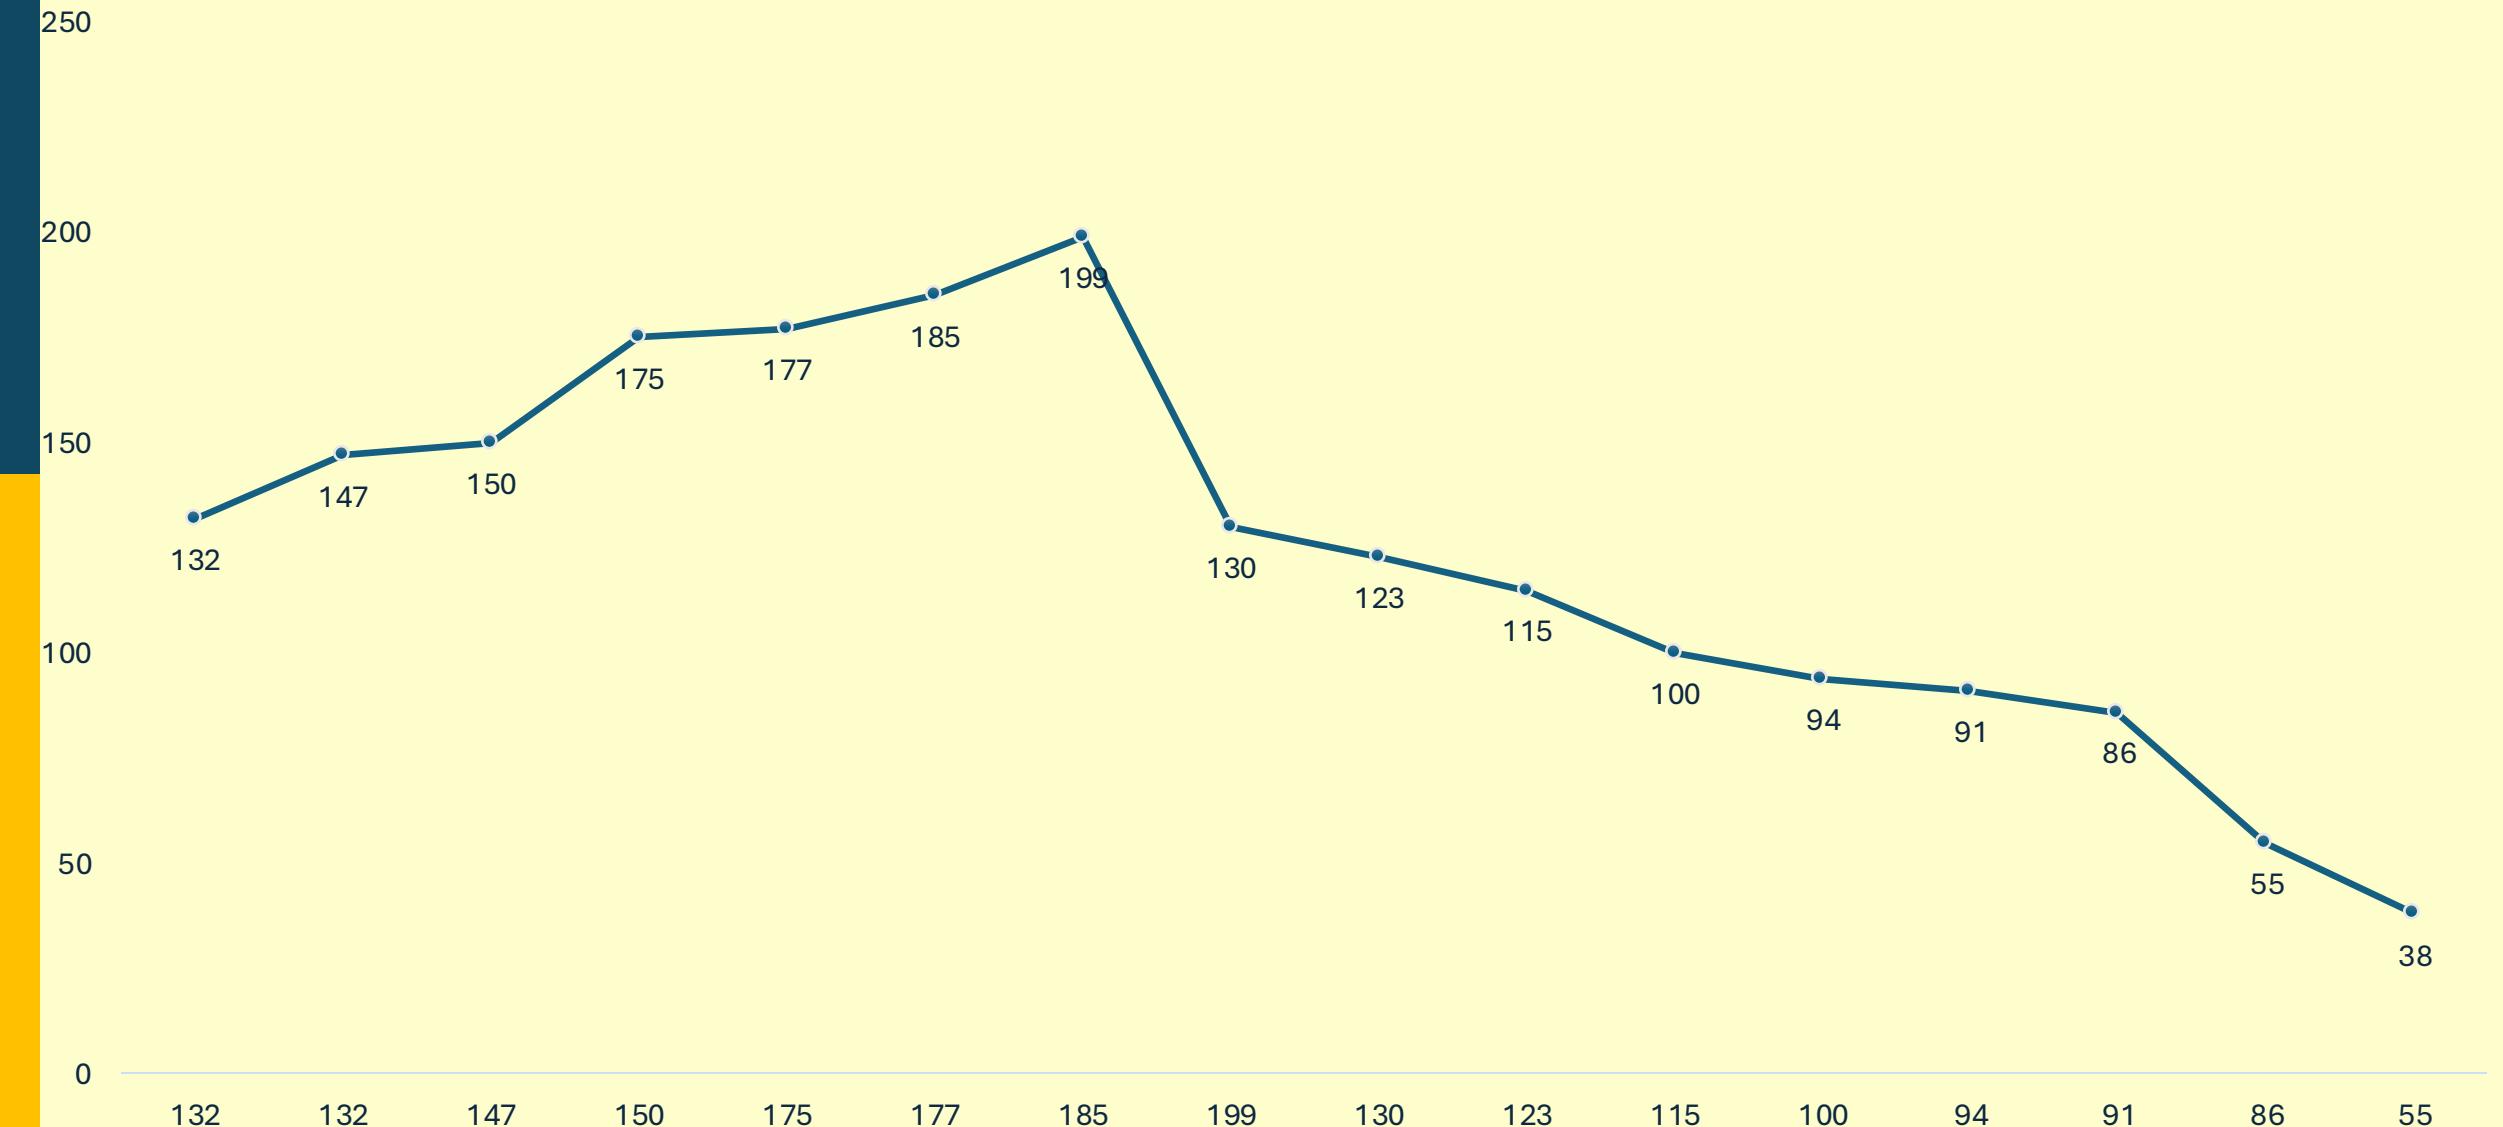
Disk Scheduling Algorithm - SCAN

| Current Position | Direction | Path Segment & Requests Serviced | Justification | Head Movement | Total Movement |
|------------------|-----------|-----------------------------------|-------------------------------------|--------------------|----------------|
| 132 | ↑ | 132 → 147 → 150 → 175 → 177 → 185 | Continue in initial direction (UP). | $ 185 - 132 = 53$ | 53 |
| 185 | ↑ | 185 → 199 | Continue to end of disk. | $ 199 - 185 = 14$ | 67 |
| 199 | ↓ | Reverse Direction | At disk end, reverse to ↓. | 0 | 67 |
| 199 | ↓ | 199 → 130 → 123 → ... → 38 | Service all remaining requests. | $ 199 - 38 = 161$ | 228 |

Sweep in one direction servicing all requests, then reverse at the disk's end and sweep back.

Disk Scheduling Algorithm - SCAN

—●— SCAN



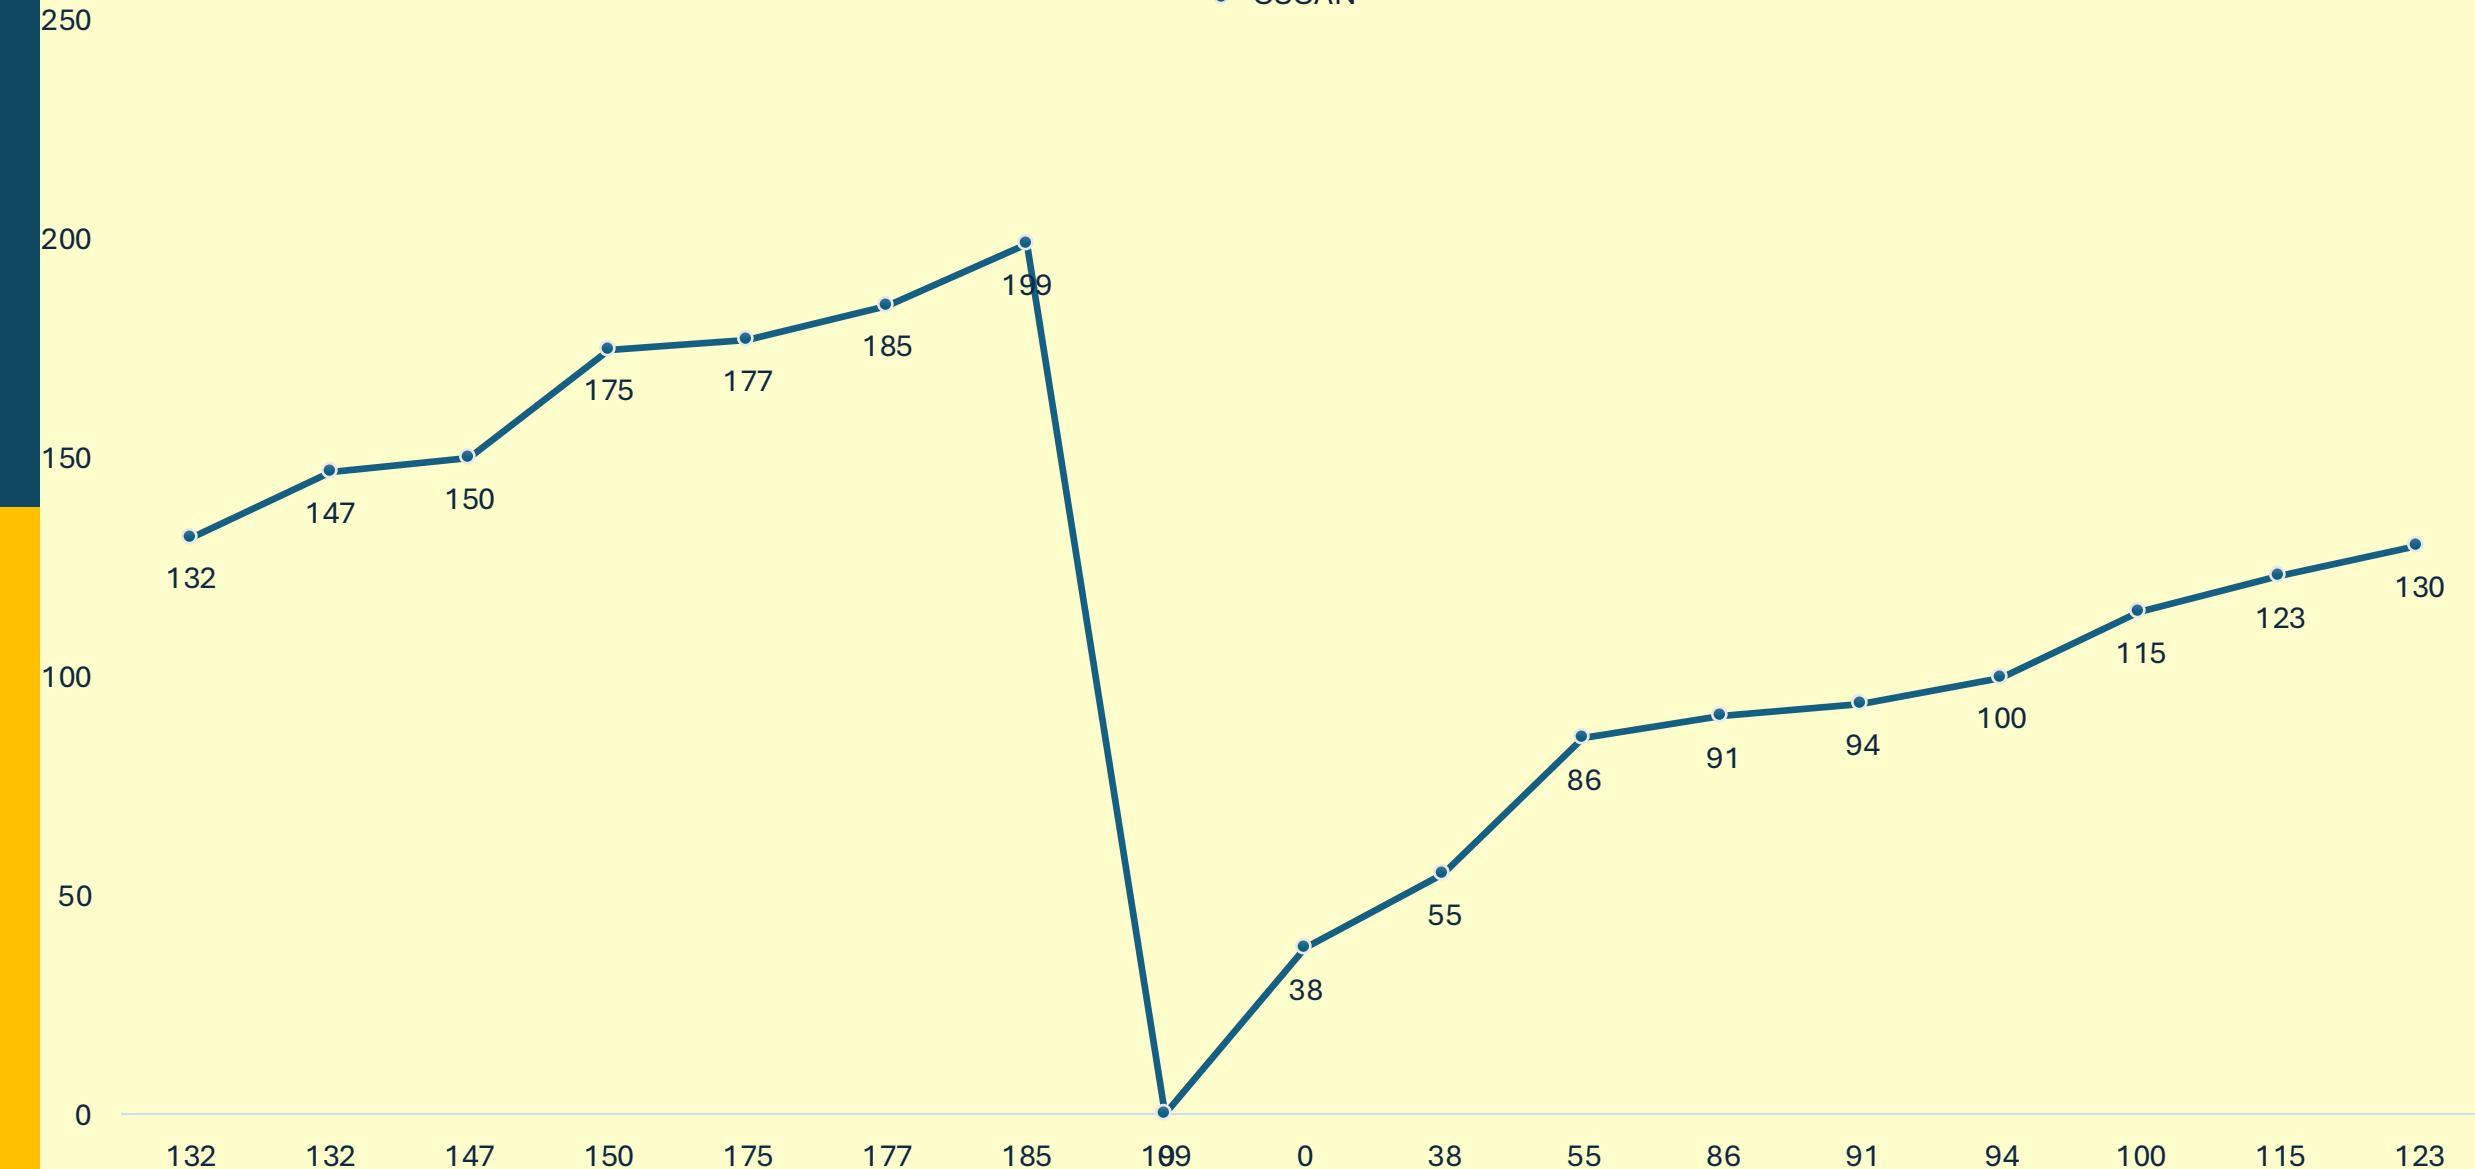
Disk Scheduling Algorithm - CSCAN

| Current Position | Direction | Path Segment & Requests Serviced | Justification | Head Movement | Total Movement |
|------------------|-----------|-----------------------------------|--|--------------------|----------------|
| 132 | ↑ | 132 → 147 → 150 → 175 → 177 → 185 | Continue in initial direction (UP). | $ 185 - 132 = 53$ | 53 |
| 185 | ↑ | 185 → 199 | Continue to end of disk. | $ 199 - 185 = 14$ | 67 |
| 199 | JUMP | 199 → 0 | At disk end, jump to start. | $ 199 - 0 = 199$ | 266 |
| 0 | ↑ | 0 → 38 → 55 → ... → 130 | Service requests from start in UP direction. | $ 130 - 0 = 130$ | 396 |

Sweep in one direction only. After reaching the end, jump back to the start and sweep again.

Disk Scheduling Algorithm - CSCAN

—●— CSCAN

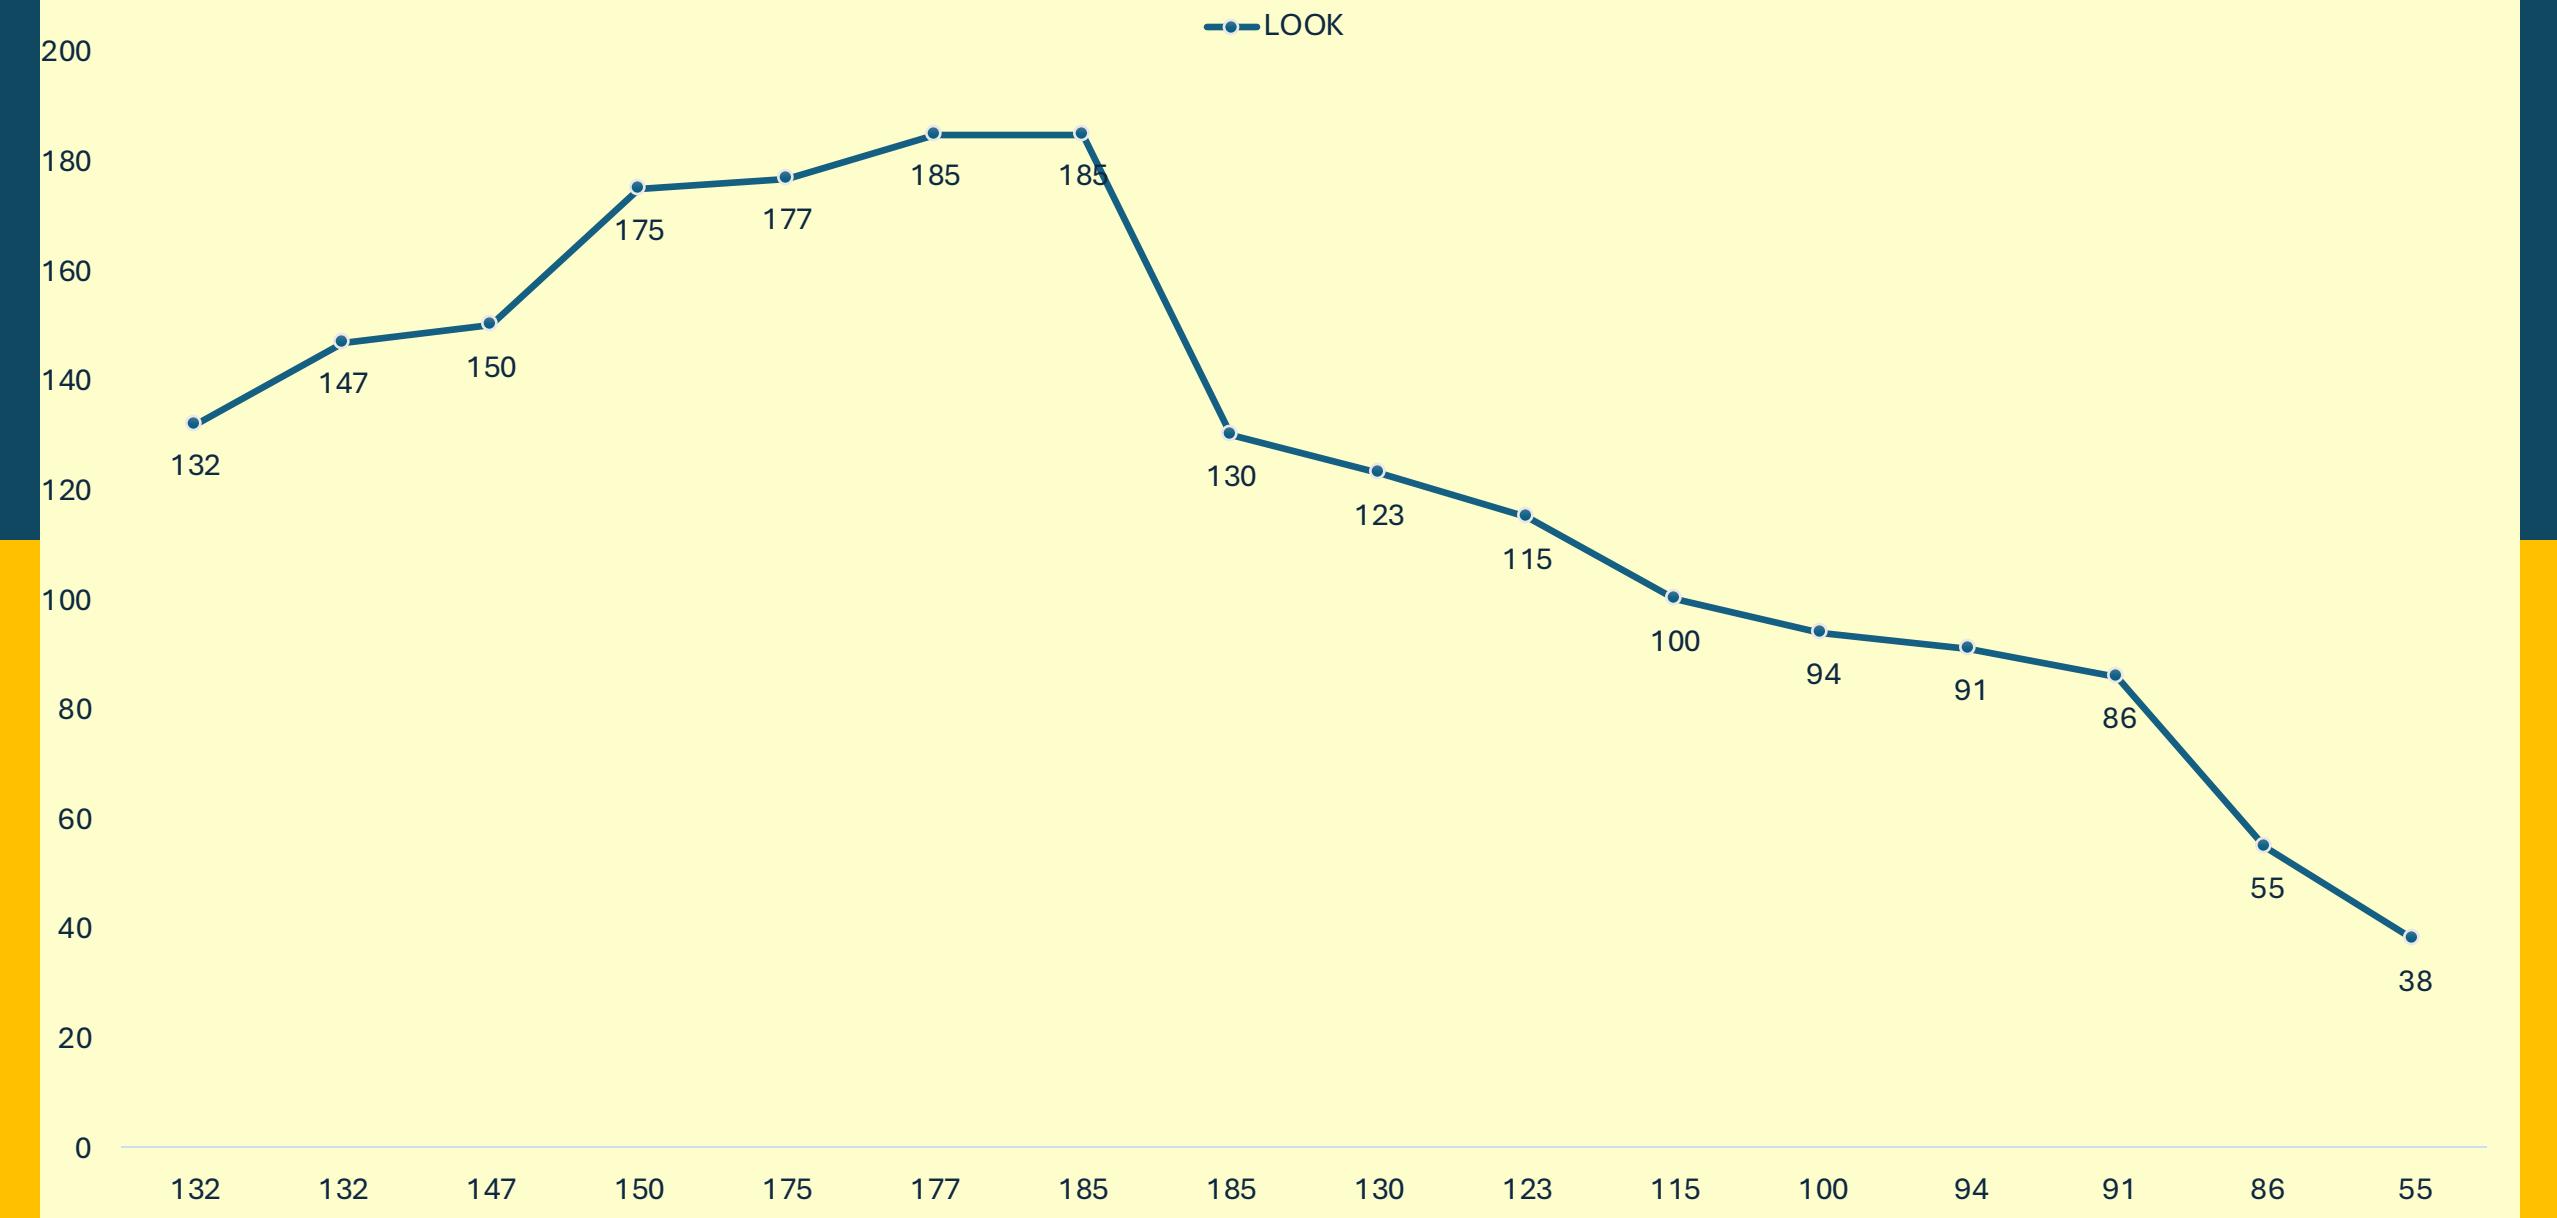


Disk Scheduling Algorithm - Look

| Current Position | Direction | Path Segment & Requests Serviced | Justification | Head Movement | Total Movement |
|------------------|-----------|-----------------------------------|--------------------------------------|--------------------|----------------|
| 132 | ↑ | 132 → 147 → 150 → 175 → 177 → 185 | Continue in initial direction (UP). | $ 185 - 132 = 53$ | 53 |
| 185 | ↓ | Reverse Direction | At last request (185), reverse to ↓. | 0 | 53 |
| 185 | ↓ | 185 → 130 → 123 → ... → 38 | Service all remaining requests. | $ 185 - 38 = 147$ | 200 |

Like SCAN, but reverse direction after the last request in that direction (**don't go to the end**).

Disk Scheduling Algorithm - Look

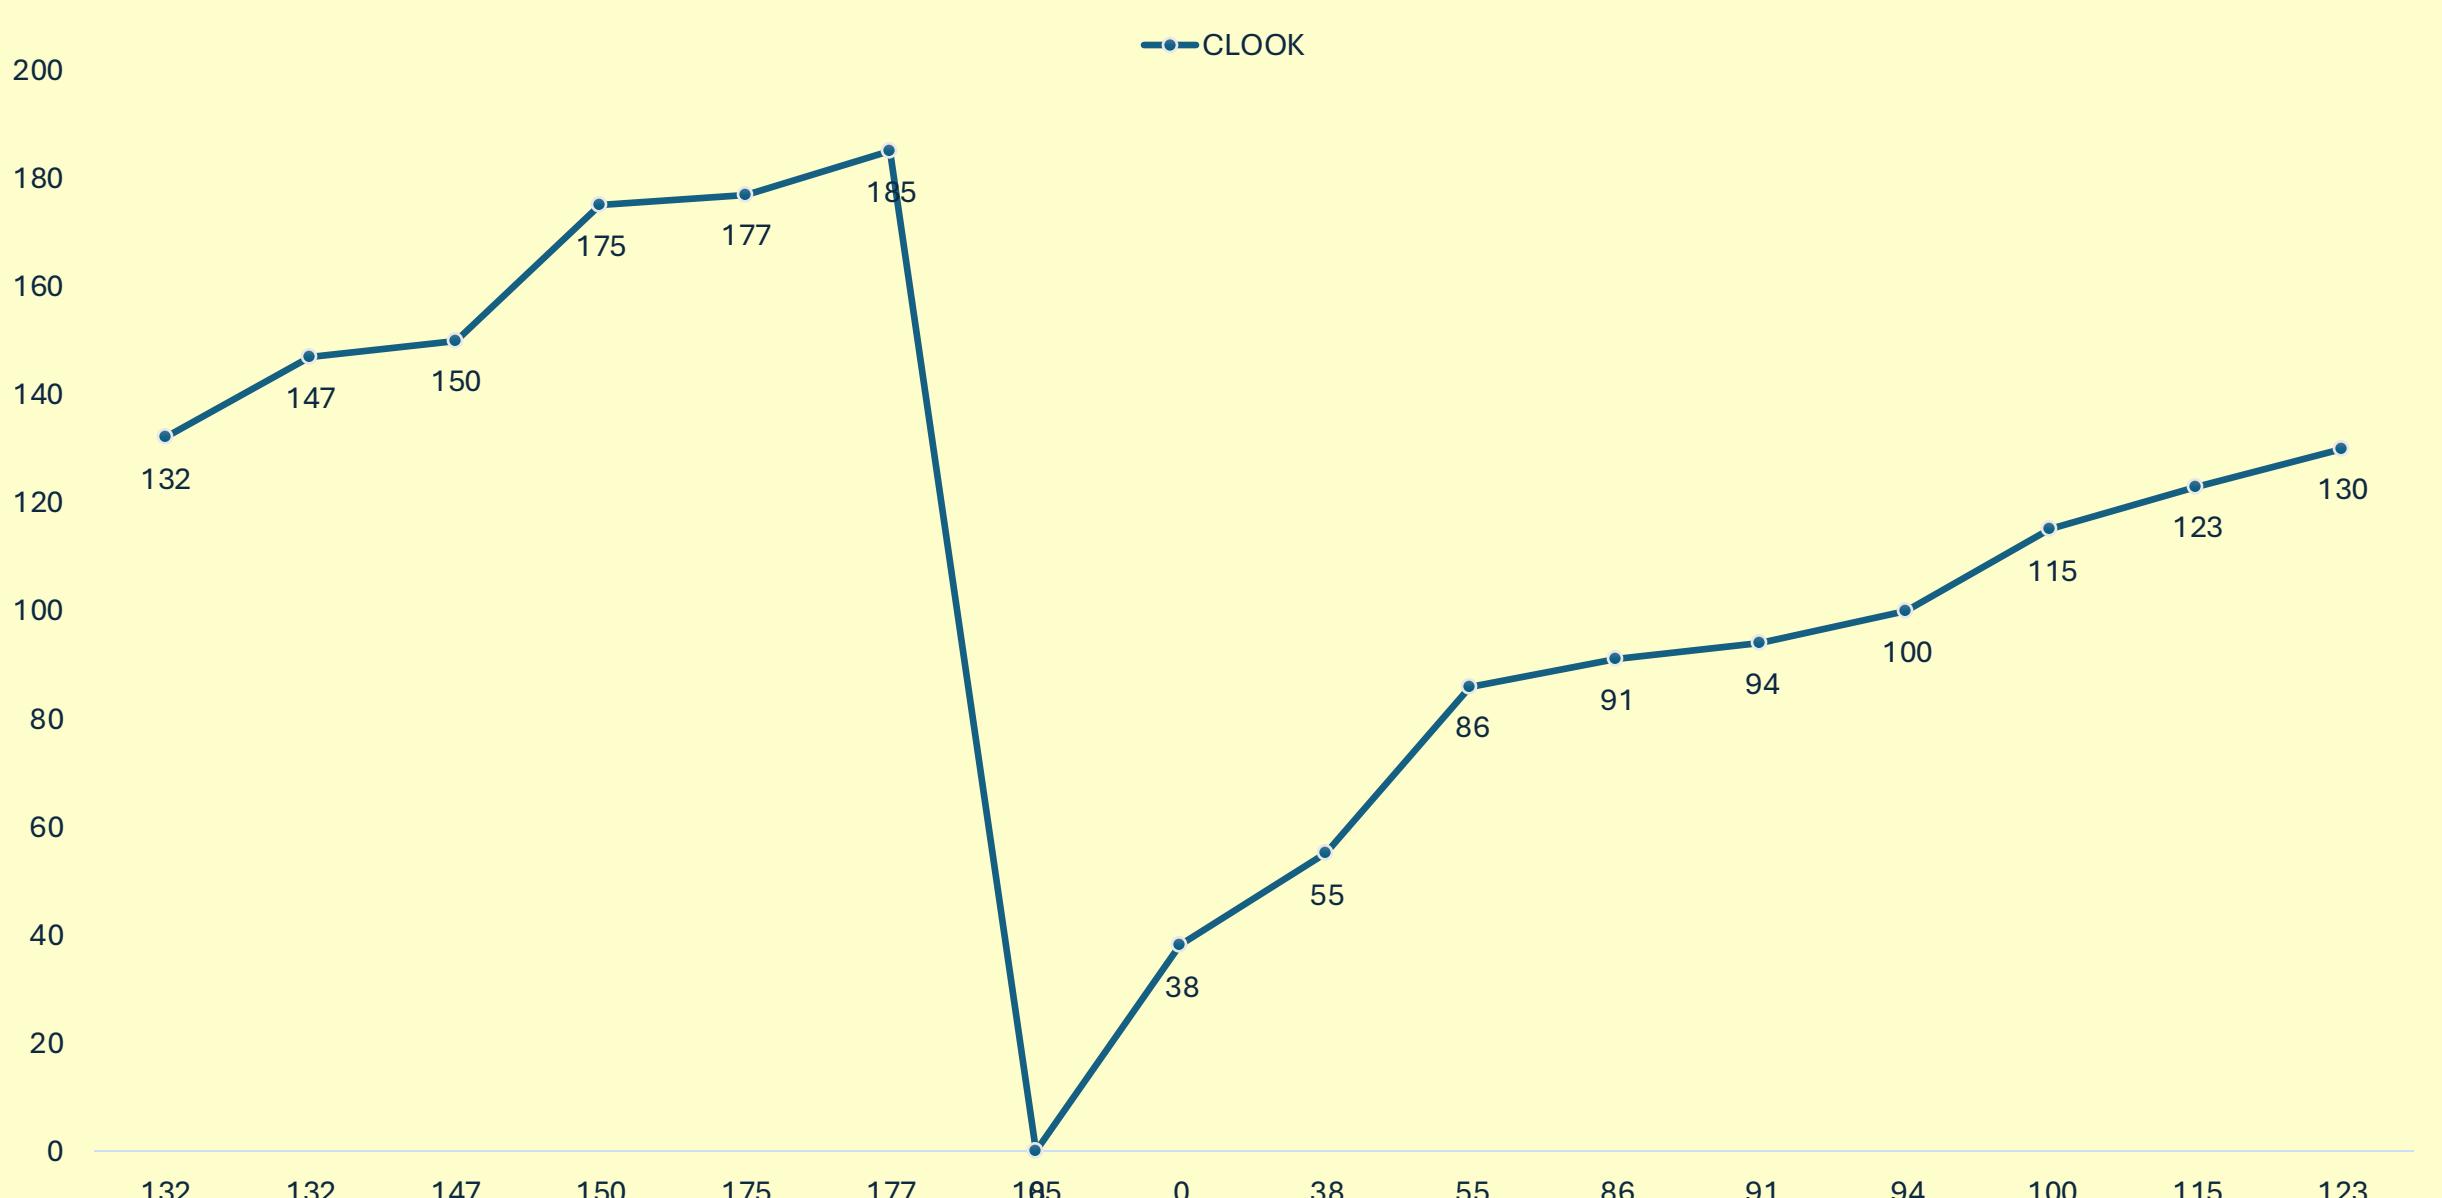


Disk Scheduling Algorithm - CLook

| Current Position | Direction | Path Segment & Requests Serviced | Justification | Head Movement | Total Movement |
|------------------|-----------|-----------------------------------|--|---------------------|----------------|
| 132 | ↑ | 132 → 147 → 150 → 175 → 177 → 185 | Continue in initial direction (UP). | $ 185 - 132 = 53$ | 53 |
| 185 | JUMP | 185 → 38 | At last request, jump to first in queue. | 0 (jump, not sweep) | 53 |
| 38 | ↑ | 38 → 55 → 86 → ... → 130 | Service remaining requests from new start. | $ 130 - 38 = 92$ | 145 |

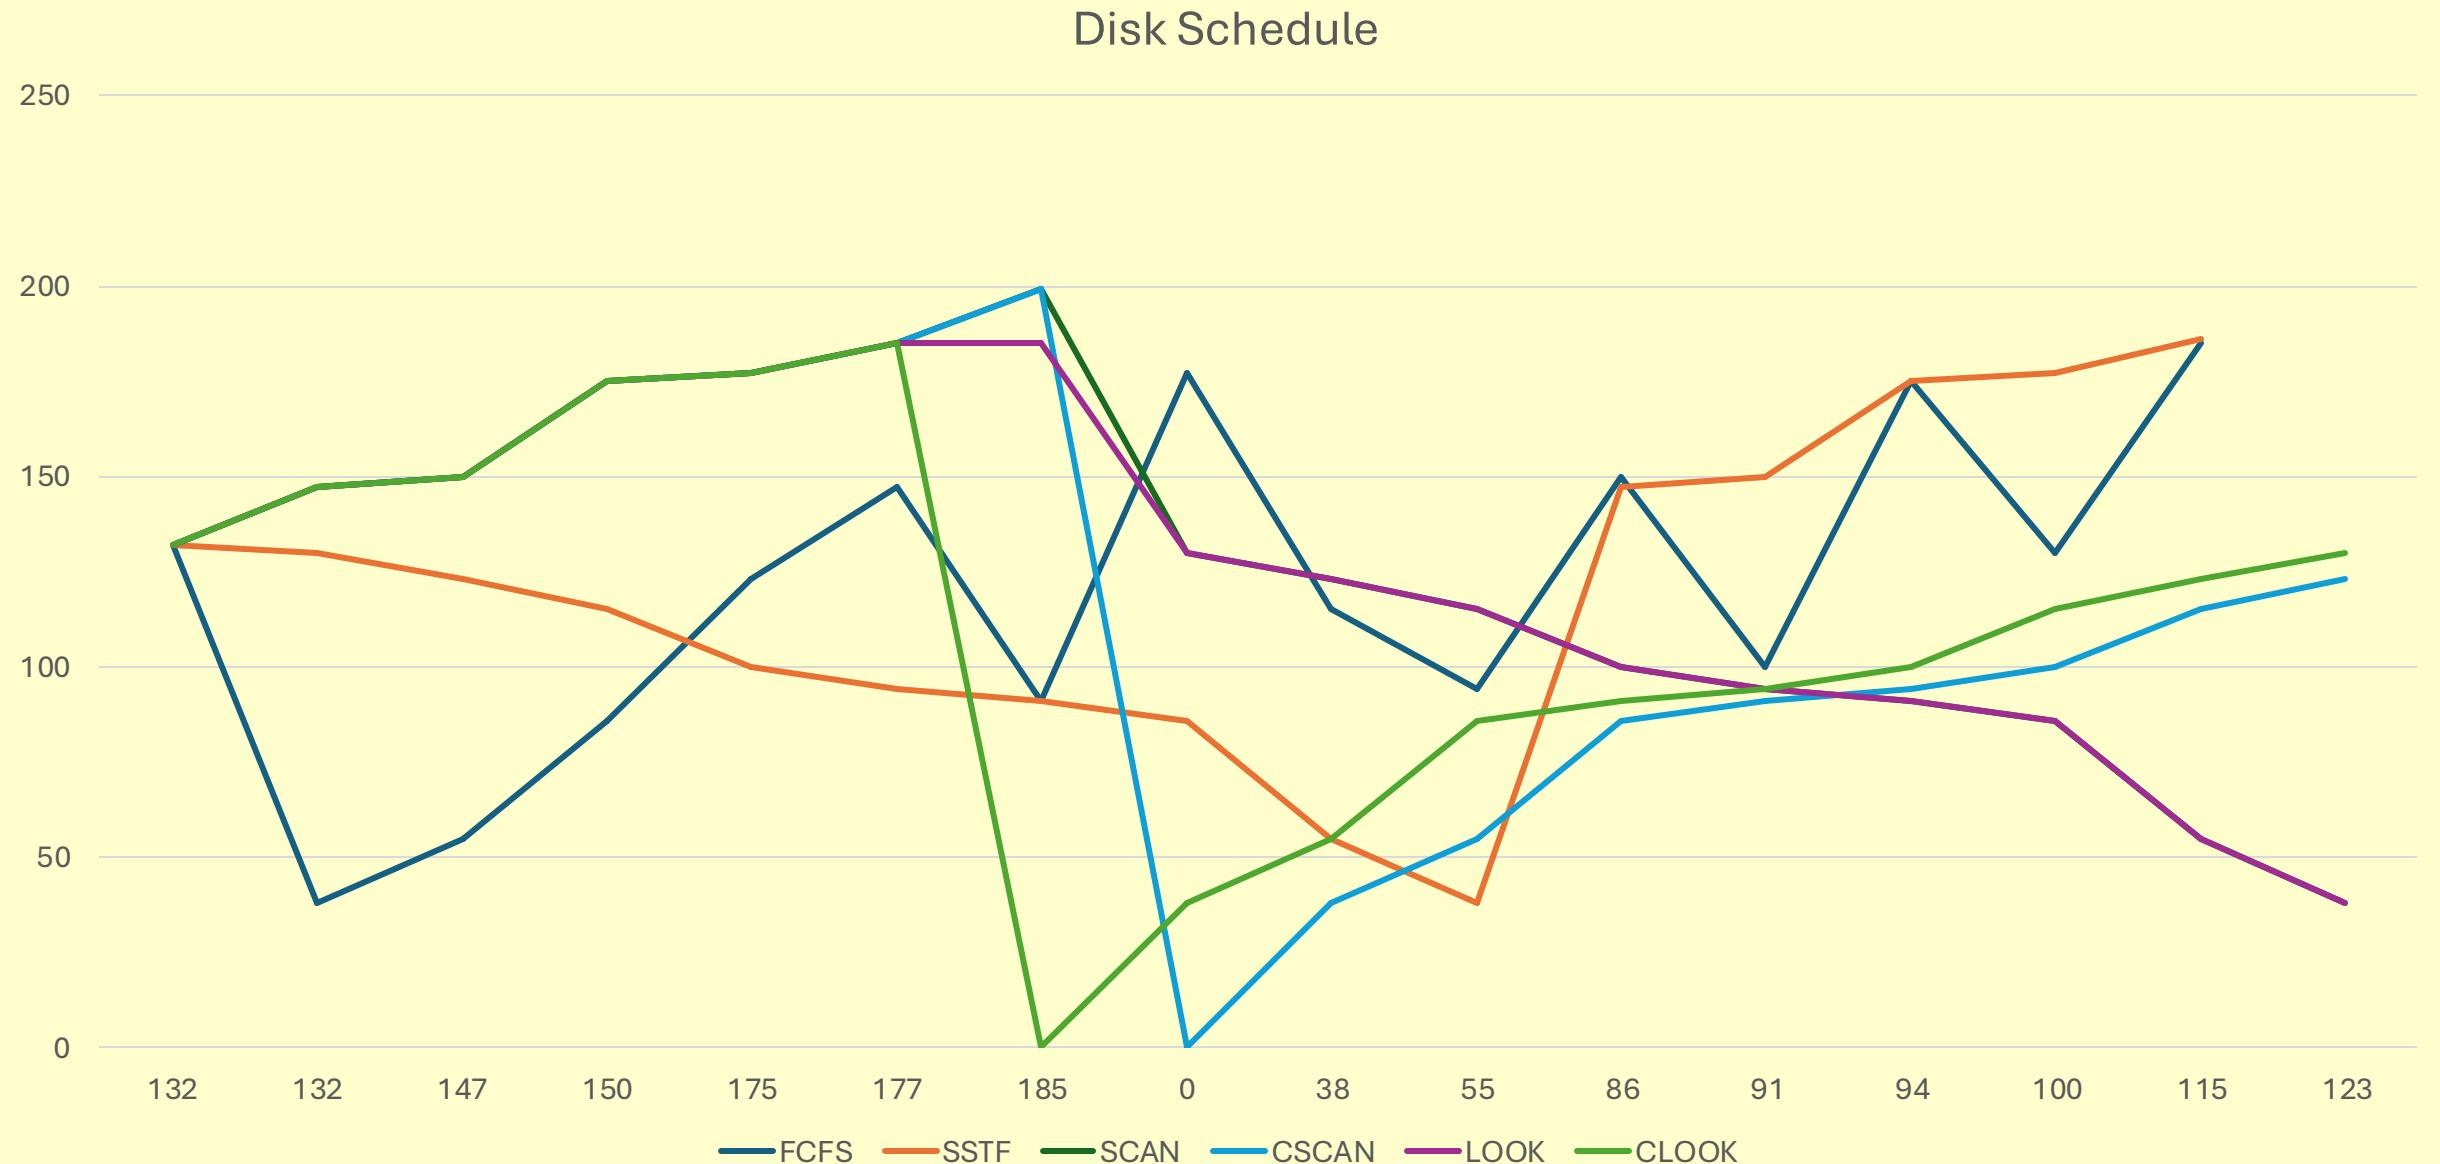
Like C-SCAN, but jump from the last request in one direction to the first request in the queue.

Disk Scheduling Algorithm - CLook



Disk Scheduling Algorithm - All

Disk Schedule



Disk Scheduling Algorithm

- ❑ Suppose the head of moving disk with 1000 tracks, numbered 0 to 999 is currently serving a request at track 132 and has just finished a request at track 120.
- ❑ If the queue of requests is kept in the FIFO order 358, 555, 861, 123, 147, 911, 177, 115, 994, 150, 100, 175, and 130, 185
- ❑ what is total head movement to satisfy these request for the following disk scheduling algorithms?
 - ❑(a) FCFS (b)LOOK (C) SCAN